

Lab 2

Rolf Sievert (rolsi701)

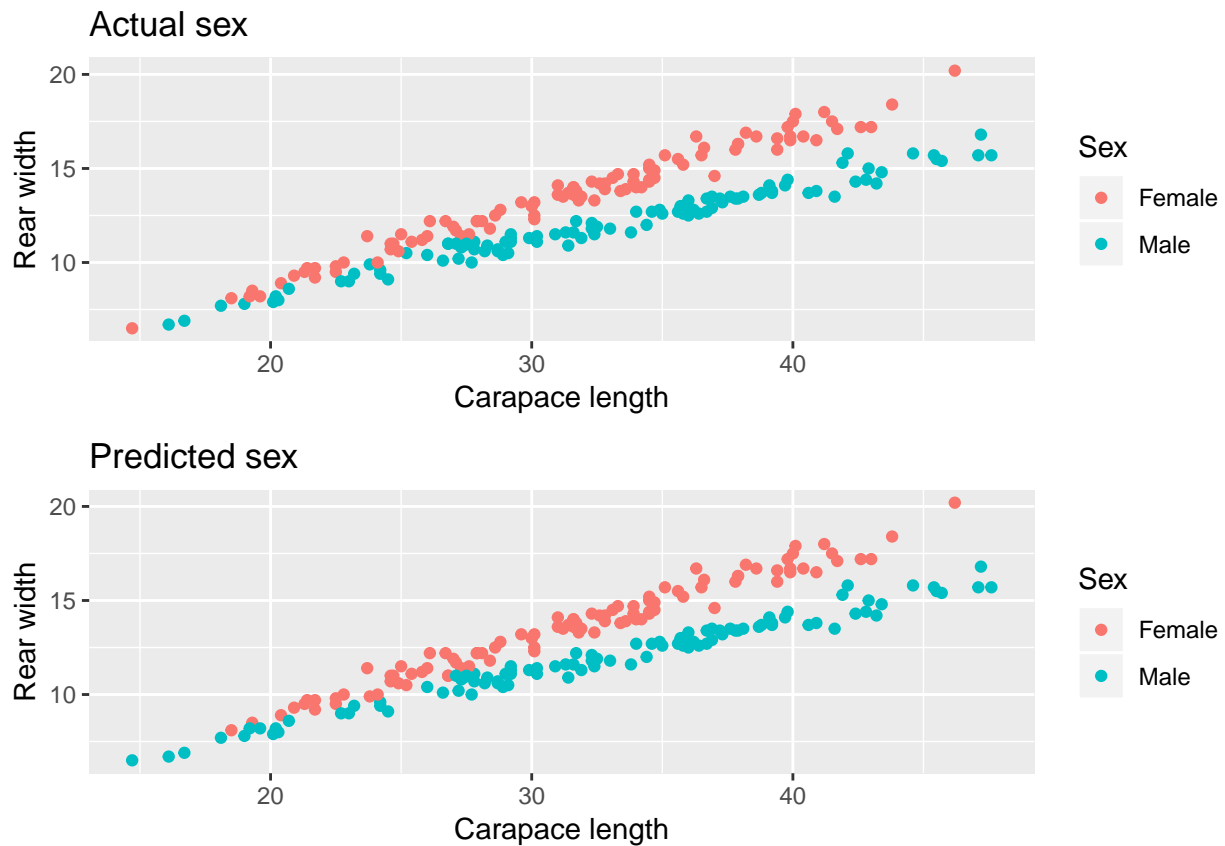
2018-12-10

Contents

Assignment 1	2
Step 1 & 2	2
Step 3	2
Step 4	3
Assignment 2	4
Step 1	4
Step 2	4
Step 3	5
Step 4	5
Step 5	6
Step 6	6
Assignment 4	7
Step 1	7
Step 2	9
Step 3	10
Appendix	12
Assignment 1	12
Assignment 2	14
Assignment 4	18

Assignment 1

Step 1 & 2



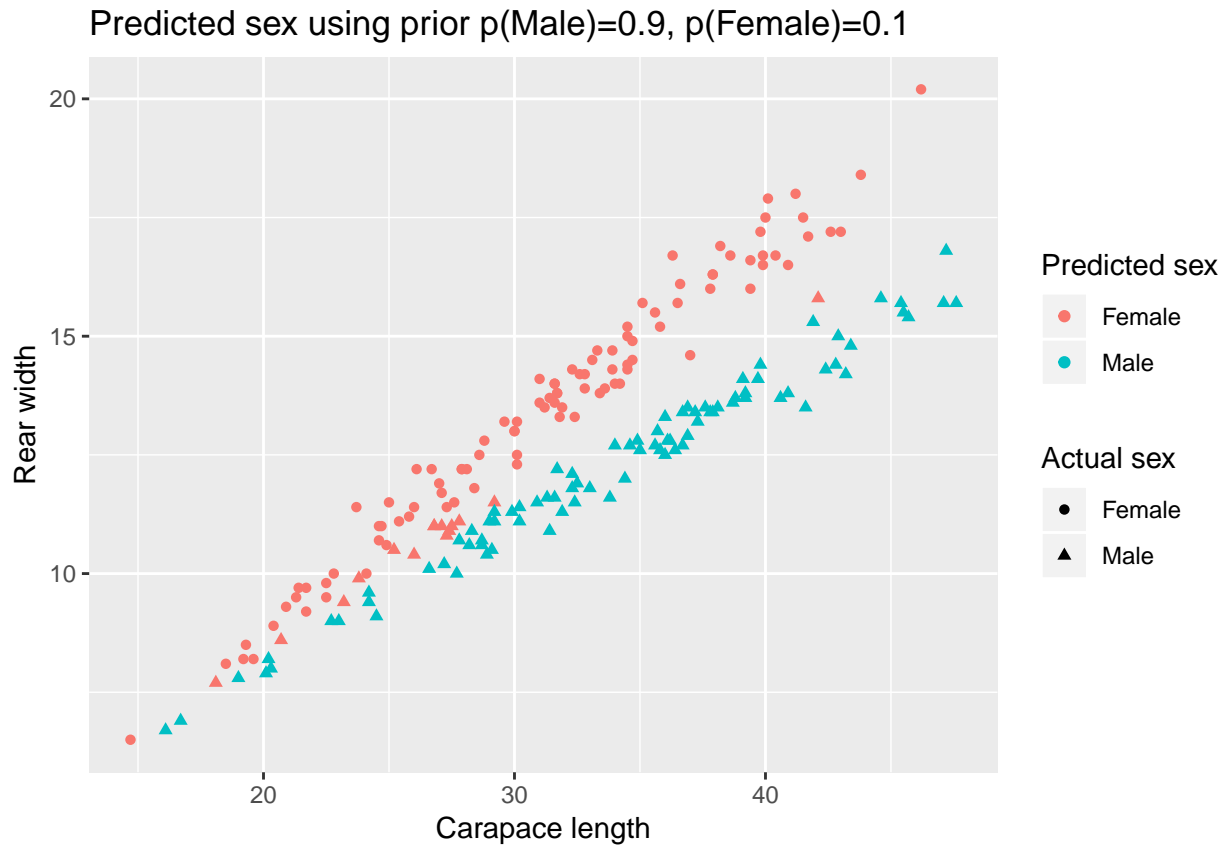
Misclassification rate for lda: 0.035

Since the sexes seems to follow two separate lines, they are separatable. Therefore it's appropriate to classify this data with a linear discriminant analysis. The only problem is where these two lines intersect and the sexes are impossible to differ. Some data may also be crossing this line and get a false classification.

The misclassification rate is really good, which strengthens the argument that the classification method is appropriate.

Step 3

This time, the same prediction is made, but with the priors $p(\text{Male}) = 0.9$ and $p(\text{Female}) = 0.1$. See plot below.



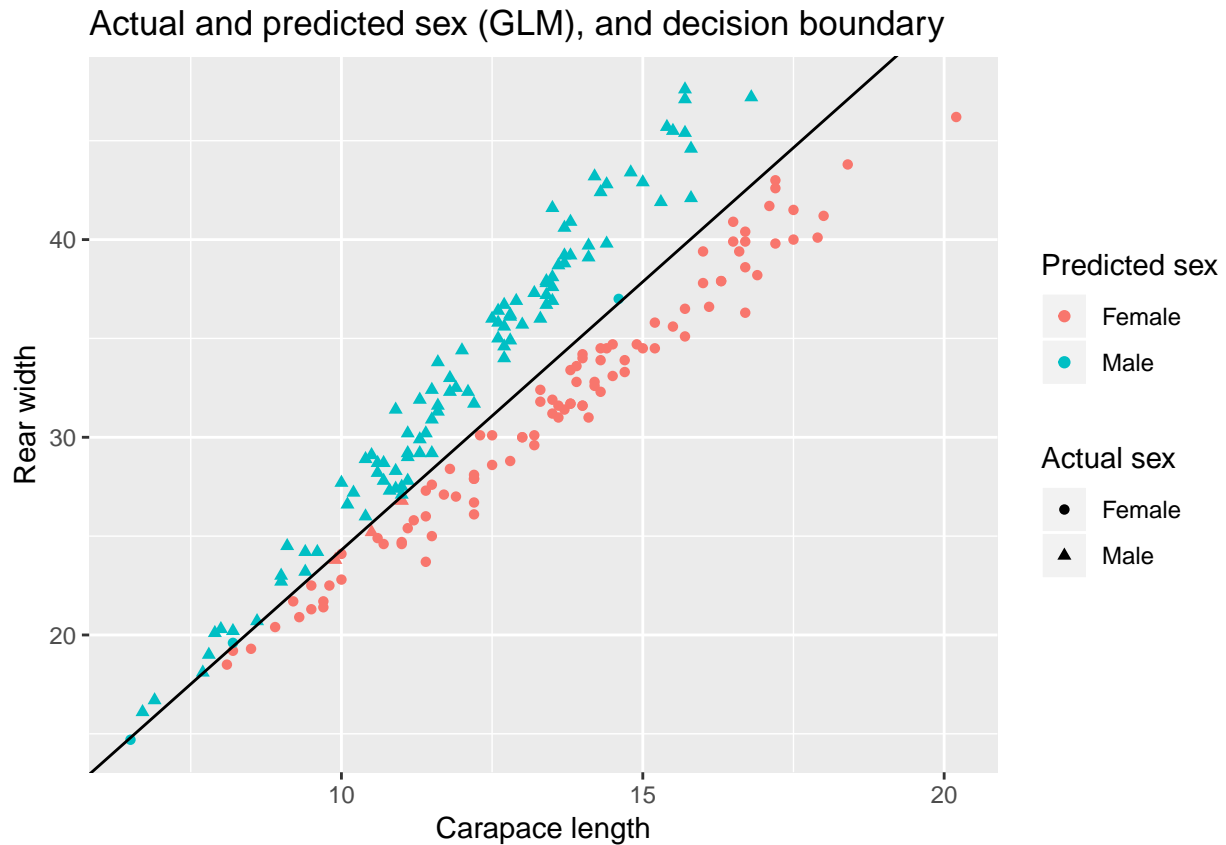
```
## Misclassification rate for lda with p(Male=0.9), p(Female)=0.1: 0.075
```

This time, the misclassification is greater. That is because the priors are not according to the actual distribution of sexes.

Step 4

Here the equation of the decision boundary is plotted together with classification using GLM.

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```
## Misclassification rate for GLM: 0.035
```

The misclassification rate is the same as in step 2.

Assignment 2

Step 1

Firstly, the data was separated into training, validation and test data.

Step 2

```
## Misclassification on deviance with test: 0.268
```

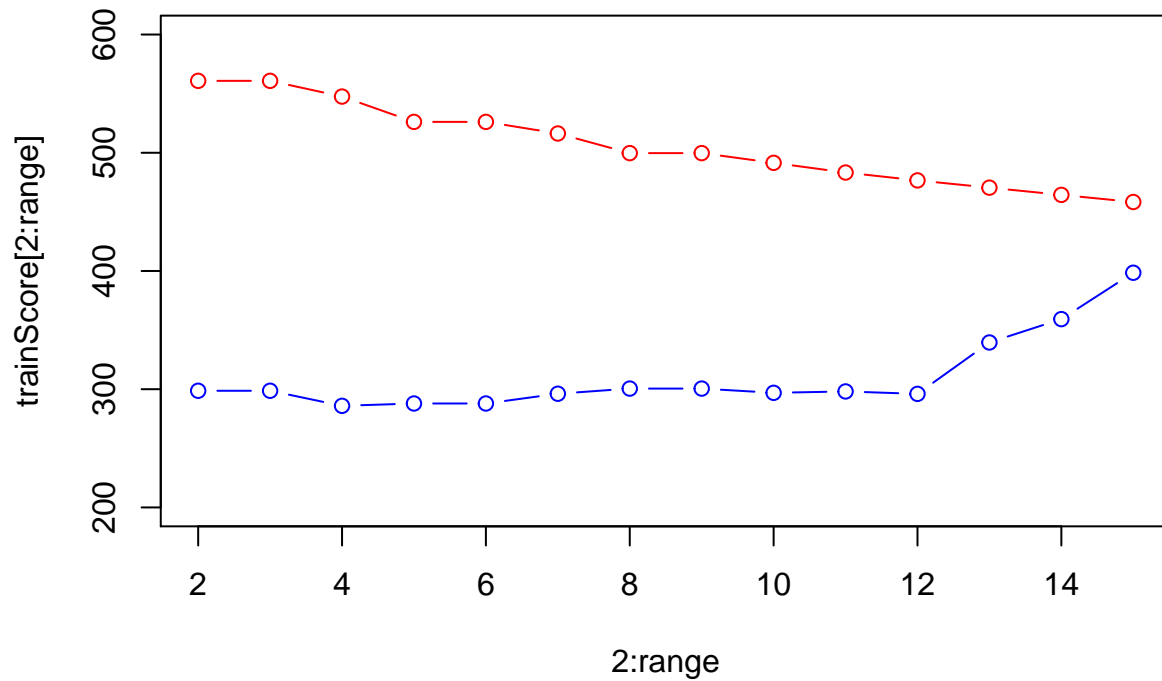
```
## Misclassification on deviance with train: 0.212
```

```
## Misclassification on gini with test: 0.368
```


```
## Misclassification on gini with train: 0.24
```

Here, misclassification rate is lower when using deviance as measure of impurity.

Step 3



```
## Misclassification on optimal tree: 0.256
## Tree depth is 3 as can be seen in the plot.
```



```
## Used variables in optimal tree:
## 'savings' 'duration' 'history'
```

Step 4

```
## Confusion table of naïve bayes (test):
##           Actual
## Predicted bad good
##      bad   46   49
##      good  30  125
```

```
## Misclassification with naïve bayes (test): 0.316
```


```
## Confusion table of naïve bayes (train):
```

```
##           Actual
## Predicted bad good
##      bad   95   98
##      good  52  255
```

```
## Misclassification with naïve bayes (train): 0.3
```

Naïve Bayes has much better result than in step 3.

Step 5



Naïve Bayes has better ratio between TPR and FPR. The only exception is around $\pi = 0.75$, as can be seen in the graph.

Step 6

Using loss matrix with naïve bayes.

```
## Confusion table of naïve bayes (using test data):
```

```
##           Actual
## Predicted bad good
##      bad   71  122
##      good   5   52
```

```
## Misclassification with naïve bayes (using test data): 0.508
```

```
## Confusion table of naïve bayes (using train data):
```

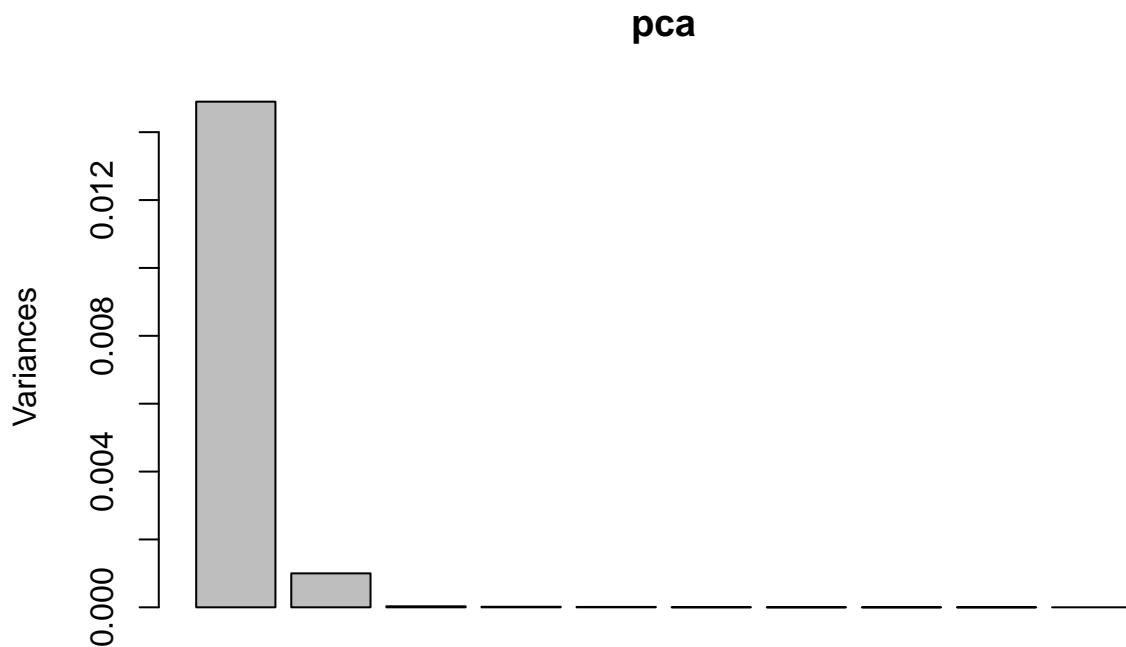
```
##           Actual
## Predicted bad good
##      bad 137 263
##      good  10  90
```

```
## Misclassification with naïve bayes (using train data): 0.546
```

The misclassification is much greater. But the confusion matrix is more favorable from an economic point of view for a company since less are predicted good that are actually bad.

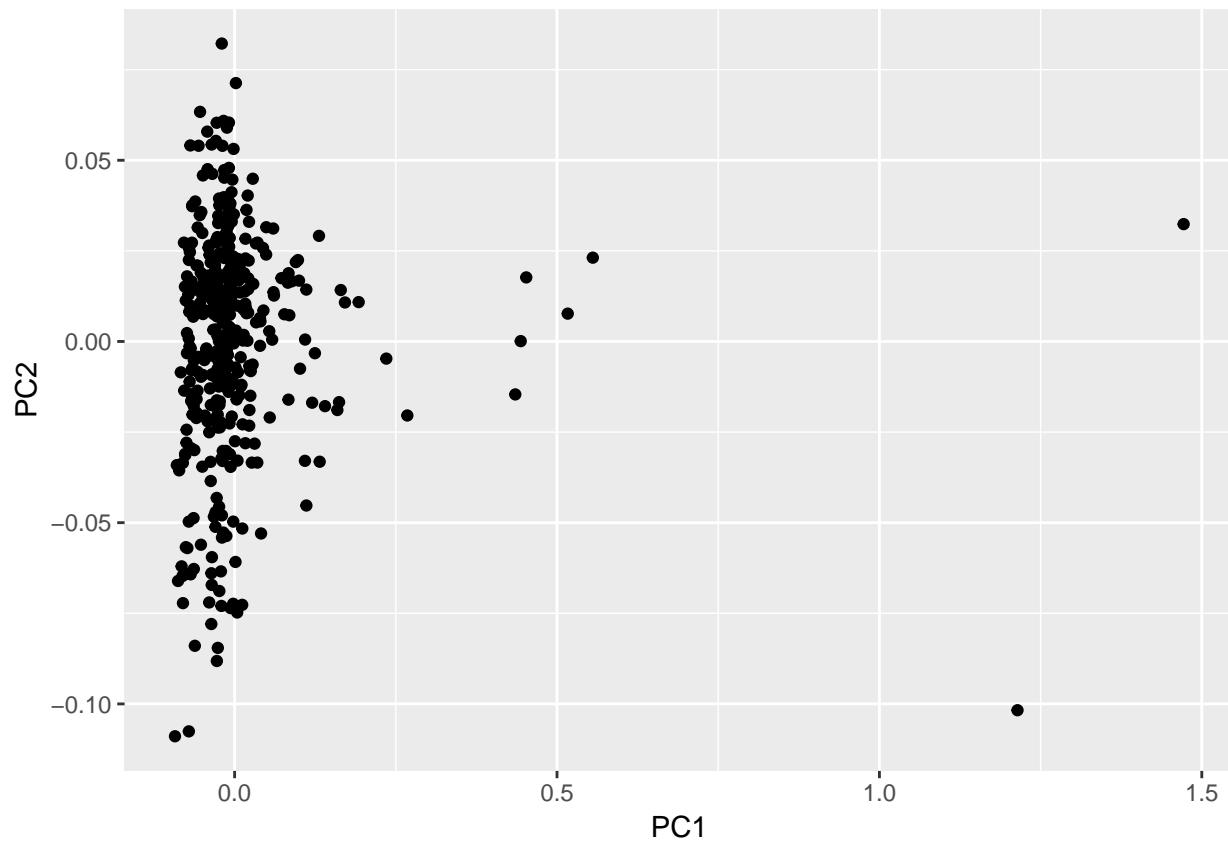
Assignment 4

Step 1



```
## [1] "93.332" "6.263" "0.185" "0.101" "0.068" "0.025" "0.009"
## [8] "0.003" "0.003" "0.002" "0.001" "0.001" "0.001" "0.001"
## [15] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [22] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [29] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [36] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [43] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [50] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [57] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [64] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [71] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [78] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [85] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [92] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [99] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
```


```
## [106] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [113] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## [120] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
## 99% of the variance can be explained by first two PC.
```




Some diesel fuels depend a lot on PC1, compared to other fuels, I would say that these are unusual.

Step 2

Traceplot, PC1



Index
Traceplot, PC2




PC2 is explained by mainly a few original features. This can be seen in the graph since it's almost zero until Index is around 100-130.

Step 3


a

```
## Centering
## Whitening
## Symmetric FastICA using logcosh approx. to neg-entropy function
## Iteration 1 tol = 0.01930239
## Iteration 2 tol = 0.01303959
## Iteration 3 tol = 0.002393582
## Iteration 4 tol = 0.0006708454
## Iteration 5 tol = 0.0001661602
## Iteration 6 tol = 3.521604e-05
```

Traceplot, PC1

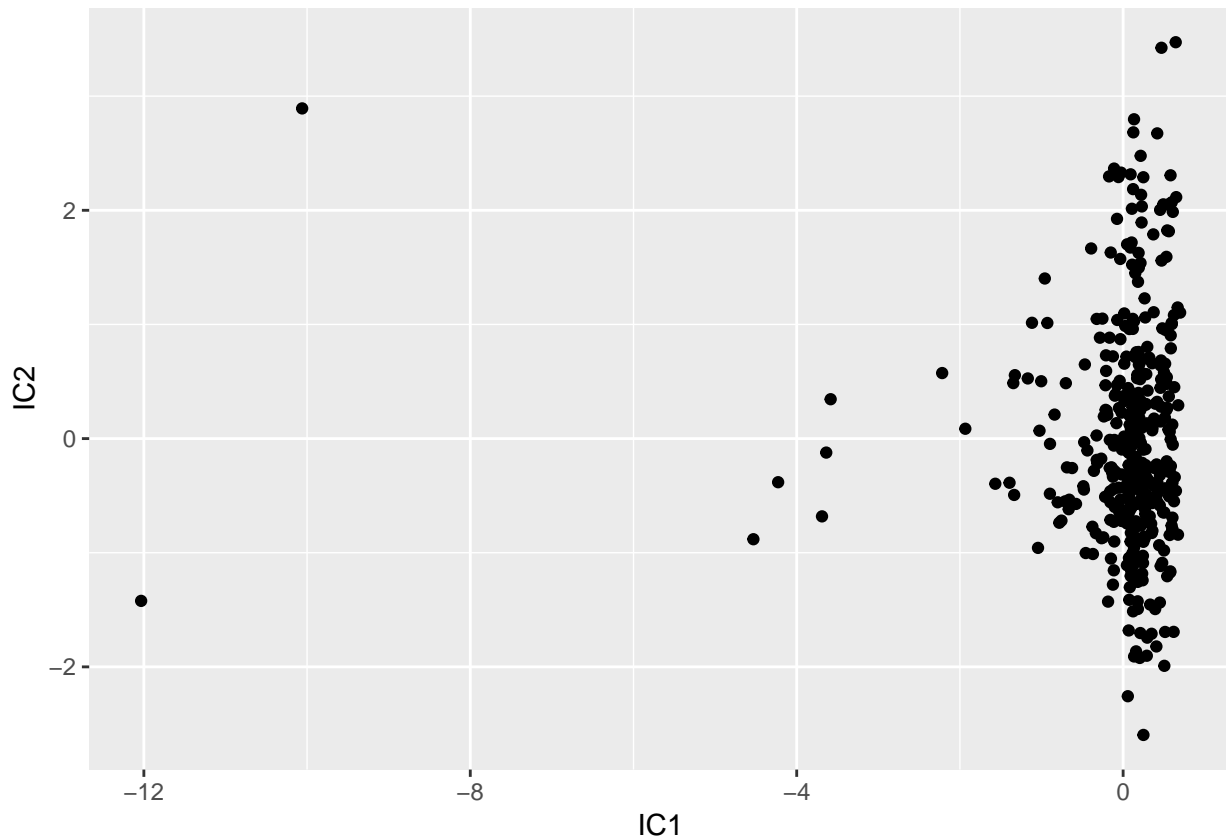


Traceplot, PC2



W' is an estimation of an un-mixing matrix.

b



This plot is very similar to the one in step 1, but it is flipped and scaled.

Appendix

Assignment 1

```
# Add libraries
library("gridExtra")
library("ggplot2")
library("MASS")

# Set working directory
setwd("~/courses/tdde01/lab2")

# Read data
crabs = read.csv("australian-crabs.csv")

crabs.plot = ggplot(data = crabs,
                    mapping = aes(x=CL,
                                  y=RW,
                                  color=sex)) +

geom_point() +
ggtitle("Actual sex") +
```

```

scale_x_continuous(name = "Carapace length") +
scale_y_continuous(name = "Rear width") +
scale_color_discrete(name = "Sex")

lda.model = lda(formula = sex~CL+RW, data = crabs)
lda.predict = predict(lda.model, crabs)
crabs.predict = ggplot(data = crabs,
                      mapping = aes(x=CL,
                                    y=RW,
                                    color=lda.predict$class)) +

geom_point() +
ggtitle("Predicted sex") +
scale_x_continuous(name = "Carapace length") +
scale_y_continuous(name = "Rear width") +
scale_color_discrete(name = "Sex")

grid.arrange(crabs.plot, crabs.predict)

# Calculate misclassification
crabs.misclass = mean(crabs$sex != lda.predict$class)
cat("Misclassification rate for lda: ", crabs.misclass)

# Step 3
prior.male = 0.9
prior.female = 0.1
crabs.predict.prior = predict(lda.model,
                             crabs,
                             prior=c(Male=prior.male,
                                       Female=prior.female)
                             )
crabs.plot.prior = ggplot(data=crabs,
                         mapping=aes(x=CL,
                                     y=RW,
                                     shape=sex,
                                     color=crabs.predict.prior$class)) +

geom_point() +
ggtitle("Predicted sex using prior p(Male)=0.9, p(Female)=0.1") +
scale_x_continuous(name = "Carapace length") +
scale_y_continuous(name = "Rear width") +
scale_color_discrete(name = "Predicted sex") +
scale_shape_discrete(name = "Actual sex")
# Plot
grid.arrange(crabs.plot.prior)
# Calculate missclassification
crabs.prior.misclass = mean(crabs$sex != crabs.predict.prior$class)
cat("Misclassification rate for lda with p(Male=0.9), p(Female)=0.1: ",
    crabs.prior.misclass)

# Step 4
glm.model = glm(formula = sex~RW+CL,
                data = crabs,
                family = 'binomial')
glm.predict = predict(object = glm.model,
                     newdata = crabs,

```

```

        type = 'response')
glm.predicted.sex = ifelse(glm.predict > 0.5, "Male", "Female")

# Calculate line from model
intercept = coef(glm.model)[1]
rw = coef(glm.model)[2]
cl = coef(glm.model)[3]
border = 0.5
k = -rw/cl
m = -(intercept - border)/cl
# Print plot and line
glm.plot = ggplot(data = crabs,
                  mapping = aes(x=RW,
                                y=CL,
                                color=glm.predicted.sex,
                                shape=crabs$sex)) +

geom_point() +
ggtitle("Actual and predicted sex (GLM), and decision boundary") +
scale_x_continuous(name = "Carapace length") +
scale_y_continuous(name = "Rear width") +
scale_color_discrete(name = "Predicted sex") +
scale_shape_discrete(name = "Actual sex") +
geom_abline(intercept = m, slope = k)
grid.arrange(glm.plot)

# Print misclassification
glm.misclass = mean(crabs$sex != glm.predicted.sex)
cat("Misclassification rate for GLM: ",
    glm.misclass)

```

Assignment 2

```

# Step 1
# Add libraries
library("gridExtra")
library("ggplot2")
library("tree")
library("MASS")
library("e1071")

# Set working directory
setwd("~/courses/tdde01/lab2")

# Read data
scores = read.csv2("creditscoring.csv")
# Read as strings
scores$good_bad = as.factor(scores$good_bad)

# Split data into train/val/test
n=dim(scores)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.5))
train=scores[id,]

```

```

id1=setdiff(1:n, id)
set.seed(12345)
id2=sample(id1, floor(n*0.25))
val=scores[id2,]
id3=setdiff(id1,id2)
test=scores[id3,]

# Step 2
# Create tree
tree.deviance = tree(formula = good_bad~.,
                      data = train,
                      split="deviance")
tree.gini = tree(formula = good_bad~.,
                 data = train,
                 split="gini")

# Make predictions
predict.deviance.test = predict(tree.deviance, test, type = "class")
predict.deviance.train = predict(tree.deviance, train, type = "class")
predict.gini.test = predict(tree.gini, test, type = "class")
predict.gini.train = predict(tree.gini, train, type = "class")

# Calculate misclassification rates
misclass = function(predicted, true) {
  return(mean(predicted != true))
}

deviance.test.misclass = misclass(predict.deviance.test, test$good_bad)
deviance.train.misclass = misclass(predict.deviance.train, train$good_bad)
gini.test.misclass = misclass(predict.gini.test, test$good_bad)
gini.train.misclass = misclass(predict.gini.train, train$good_bad)

# Print misclassification rates
cat("Misclassification on deviance with test: ",
    deviance.test.misclass)
cat("\nMisclassification on deviance with train: ",
    deviance.train.misclass)
cat("\nMisclassification on gini with test: ",
    gini.test.misclass)
cat("\nMisclassification on gini with train: ",
    gini.train.misclass, "\n")

# Step 3
# Lists of scores
range = 15
pruned=rep(0, range)
trainScore=rep(0, range)
testScore=rep(0, range)

test.tree = tree(formula=good_bad~., data=train)
for(i in 2:range) {
  # Prune the tree
  prunedTree=prune.tree(test.tree, best=i)
  # Make prediction on validation data

```

```

    pred=predict(prunedTree, newdata=val, type="tree")
    # Append scores
    trainScore[i]=deviance(prunedTree)
    testScore[i]=deviance(pred)
}

# Plot the scores
plot(2:range, trainScore[2:range], type="b", col="red", ylim=c(200, 600))
points(2:range, testScore[2:range], type="b", col="blue")

optLeaves = 4

# Info on optimal tree
optTree = prune.tree(tree.deviance, best=optLeaves)
summary(optTree)
optTree.predict = predict(optTree, newdata=test)
optTree.predict.string = ifelse(optTree.predict[2] > 0.5, "good", "bad")
optTree.misclass = misclass(optTree.predict.string, test$good_bad)
cat("\nMisclassification on optimal tree: ",
    optTree.misclass)

cat("\nTree depth is 5 as can be seen in the plot. \n")
plot(optTree)
cat("\n\nUsed variables in optimal tree: \n'savings' 'duration' 'history' 'age' 'purpose'\n\n")

# Step 4
# Create model and classify
model.bayes = naiveBayes(formula = good_bad~., data=train)
predict.bayes.test = predict(model.bayes, newdata=test, type="class")
predict.bayes.train = predict(model.bayes, newdata=train, type="class")

# Print info on classification with prediction on test
table.bayes = table(predict.bayes.test, test$good_bad)
cat("Confusion table of naïve bayes:")
print(table.bayes)

misclass.bayes = mean(predict.bayes.test != test$good_bad)
cat("Misclassification with naïve bayes: ", misclass.bayes, "\n\n")

# Print info on classification with prediction on test
table.bayes = table(predict.bayes.train, train$good_bad)
cat("Confusion table of naïve bayes:")
print(table.bayes)

misclass.bayes = mean(predict.bayes.train != train$good_bad)
cat("Misclassification with naïve bayes: ", misclass.bayes, "\n\n")

# Step 5
# Calculate TPR and FPR of optimal tree and bayes
getROC = function(pred, pi) {
  tpr = c()
  fpr = c()
  for (p in pi) {
    # Change probabilities to strings

```



```

    tmp = ifelse(pred[, 'good'] > p, "good", "bad")
    # Get confusion matrix
    cm = table(predicted=tmp, actual=test$good_bad)
    if('good' %in% rownames(cm)) {
      # Calculate TPR, first dim of cm is predicted
      t = cm['good', 'good'] / sum(cm[, 'good'])
      # Calculate FPR
      f = cm['good', 'bad'] / sum(cm[, 'bad'])
      # Append to list of values
      tpr = c(tpr, ifelse(is.finite(t), t, 0))
      fpr = c(fpr, ifelse(is.finite(f), f, 0))
    } else {
      tpr = c(tpr, 0)
      fpr = c(fpr, 0)
    }
  }
  df = data.frame(tpr, fpr)
  return(df)
}

pi = seq(0.05, 0.95, 0.05)
pred.optTree = predict(optTree, newdata=test)
pred.bayes = predict(model.bayes, newdata=test, type='raw')
opt.roc = getROC(pred.optTree, pi)
bayes.roc = getROC(pred.bayes, pi)

#Plot
roc.plot = ggplot(mapping=aes(col=Classifier)) +
  geom_point(data=opt.roc, aes(x=fpr, y=tpr, col="Optimal tree")) +
  geom_line(data=opt.roc, aes(x=fpr, y=tpr, col="Optimal tree")) +
  geom_point(data=bayes.roc, aes(x=fpr, y=tpr, col="Bayes")) +
  geom_line(data=bayes.roc, aes(x=fpr, y=tpr, col="Bayes"))

grid.arrange(roc.plot)

# Step 6
# Create model and classify
model.bayes = naiveBayes(formula = good_bad~., data=train)
predict.bayes.test = predict(model.bayes, newdata=test, type="raw")
predict.bayes.train = predict(model.bayes, newdata=train, type="raw")

# Print info on classification with prediction on test
loss = 10/1
predict.bayes.test.loss = ifelse(predict.bayes.test[, 'good'] /
                                predict.bayes.test[, 'bad'] > loss,
                                'good',
                                'bad')

table.bayes = table(Predict=predict.bayes.test.loss, Actual=test$good_bad)
cat("Confusion table of naïve bayes (using test data):")
print(table.bayes)

misclass.bayes = mean(predict.bayes.test.loss != test$good_bad)
cat("Misclassification with naïve bayes (using test data): ", misclass.bayes, "\n")

```

```

# Print info on classification with prediction on train
predict.bayes.train.loss = ifelse(predict.bayes.train[, 'good'] /
                                predict.bayes.train[, 'bad'] > loss,
                                'good',
                                'bad')

table.bayes = table(Predict=predict.bayes.train.loss, Actual=train$good_bad)
cat("Confusion table of naïve bayes (using train data): ")
print(table.bayes)

misclass.bayes = mean(predict.bayes.train.loss != train$good_bad)
cat("Misclassification with naïve bayes (using train data): ", misclass.bayes, "\n")

```

Assignment 4

```

# Step 1
# Add libraries
library("gridExtra")
library("ggplot2")
library("MASS")
library("stats")
library("fastICA")

# Set working directory
setwd("~/courses/tdde01/lab2")

# Read data
spectra = data.frame(read.csv2("NIRSpectra.csv"))
spectra = subset(spectra, select=-Viscosity)

# Calculate principal components
pca = prcomp(spectra)

# Get eigen values
lambda = pca$sdev^2
# Plot variance of each component
screplot(pca)

# Print proportion of variation
sprintf("%.2.3f", lambda/sum(lambda)*100)
cat("99% of the variance can be explained by first two PC.\n\n")

# Extract principal components
PC1 = pca$x[,1]
PC2 = pca$x[,2]

# Plot the data projected onto PC 1 and 2
plot.projection = ggplot() +
  geom_point(data=spectra, mapping=aes(x=PC1, y=PC2))
grid.arrange(plot.projection)

# Step 2
plot(pca$rotation[,1], main="Traceplot, PC1")
plot(pca$rotation[,2], main="Traceplot, PC2")

```

```

# Step 3a
set.seed(12345)
res = fastICA(X=spectra,
  n.comp=2,
  alg.typ= "parallel",
  fun = "logcosh",
  alpha = 1,
  method = "R",
  row.norm = FALSE,
  maxit= 200,
  tol = 0.0001,
  verbose = TRUE)

W.prime = res$K %*% res$W
# Traceplots of fastICA
plot(W.prime[,1], main="Traceplot, PC1")
plot(W.prime[,2], main="Traceplot, PC2")

# Step 3b
# Extract ICs
IC1 = res$S[,1]
IC2 = res$S[,2]

# Plot the data projected onto PC 1 and 2
plot.projection = ggplot() +
  geom_point(data=spectra, mapping=aes(x=IC1, y=IC2))
grid.arrange(plot.projection)

```